

# Elevating R to Supercomputers

Drew Schmidt

National Institute for Computational Sciences  
University of Tennessee, Knoxville

November 11, 2013



# The pbdR Core Team

Wei-Chen Chen<sup>1</sup>

George Ostrouchov<sup>1,2</sup>

Pragneshkumar Patel<sup>2</sup>

Drew Schmidt<sup>2</sup>



## Support

This work used resources of the [Oak Ridge Leadership Computing Facility](#) at the Oak Ridge National Laboratory, which is supported by the Office of Science of the U.S. Department of Energy under Contract No. DE-AC05-00OR22725. This work also used resources of [National Institute for Computational Sciences](#) at the University of Tennessee, Knoxville, which is supported by the Office of Cyberinfrastructure of the U.S. National Science Foundation under Award No. ARRA-NSF-OCI-0906324 for NICS-RDAV center. This work used resources of the Newton HPC Program at the University of Tennessee, Knoxville.

<sup>1</sup>Oak Ridge National Laboratory. Supported in part by the project “Visual Data Exploration and Analysis of Ultra-large Climate Data” funded by U.S. DOE Office of Science under Contract No. DE-AC05-00OR22725.

<sup>2</sup>University of Tennessee. Supported in part by the project “NICS Remote Data Analysis and Visualization Center” funded by the Office of Cyberinfrastructure of the U.S. National Science Foundation under Award No. ARRA-NSF-OCI-0906324 for NICS-RDAV center.

# Contents

1 Introduction

2 Benchmarks

3 Challenges

## Why R?

- 1 Because.
- 2 R community has growing data size problem.
- 3 HPC community has growing need for data analytics.

## Elevating R to Supercomputers

- 1 Existing code.
- 2 Syntax.
- 3 Philosophy.

## Programming with Big Data in R (pbdR)

*Productivity, Portability, Performance*





- *Free<sup>a</sup>* R packages.
- Bridging high-performance C with high-productivity of R
- Distributed data details implicitly managed.
- Methods have syntax *identical* to R.

---

<sup>a</sup>MPL, BSD, and GPL licensed

## pbdR Packages



## pbdMPI vs Rmpi: API

### Reduction Operations

#### Rmpi

```
1 # int
2 mpi.allreduce(x, type=1)
3 # double
4 mpi.allreduce(x, type=2)
```

#### pbdMPI

```
1 allreduce(x)
```

### Types in R

```
1 > is.integer(1)
2 [1] FALSE
3 > is.integer(2)
4 [1] FALSE
5 > is.integer(1:2)
6 [1] TRUE
```



## pbdMPI vs Rmpi: Performance

**Table:** Runtimes (seconds) for  $10,000 \times 10,000$  allgather with **Rmpi** and **pbdMPI**.

Cores	Rmpi	pbdMPI	Speedup
32	24.6	6.7	3.67
64	25.2	7.1	3.55
128	22.3	7.2	3.10
256	22.4	7.1	3.15

## pbdR Example Syntax

```
1 x <- x[-1, 2:5]
2 x <- log(abs(x) + 1)
3 xtx <- t(x) %*% x
4 ans <- svd(solve(xtx))
```


Look familiar?

*The above runs on 1 core with R or 10,000 cores with pbdR*

## Shared and Distributed Memory Machines


### Shared Memory

Direct access to read/change memory (one node)



### Distributed

No direct access to read/change memory.



## Shared and Distributed Memory Machines

### Shared Memory Machines

Thousands of cores



*Nautilus*, University of Tennessee  
1024 cores  
4 TB RAM

### Distributed Memory Machines

Hundreds of thousands of cores



*Kraken*, University of Tennessee  
112,896 cores  
147 TB RAM

# Contents

- 1 Introduction
- 2 Benchmarks
- 3 Challenges

## Non-Optimal Choices Throughout

- 1 Only libre software used (no MKL, ACML, etc.).
- 2 1 core = 1 MPI process.
- 3 No tuning for data distribution.

## Benchmark Data


- 1 Random normal  $N(100, 10000)$ .
- 2 Local problem size of  $\approx 43.4 \text{ MiB}$ .
- 3 Three sets: 500, 1000, and 2000 columns.
- 4 Several runs at different core sizes within each set.

## Covariance Code

```
1 x <- ddmatrix("rnorm", nrow=n, ncol=p, mean=mean, sd=sd)
2
3 cov.x <- cov(x)
```




cov()



## Linear Model Code

```
1 x <- ddmatrix("rnorm", nrow=n, ncol=p, mean=mean, sd=sd)
2 beta_true <- ddmatrix("runif", nrow=p, ncol=1)
3
4 y <- x %*% beta_true
5
6 beta_est <- lm.fit(x=x, y=y)$coefficients
```


`lm.fit()`

But wait! There's more...

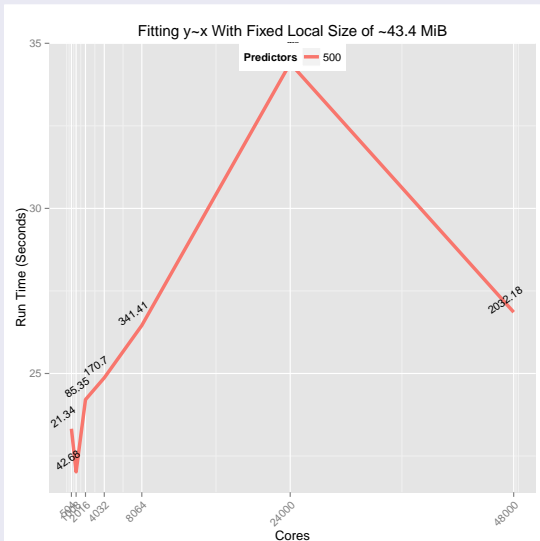
*Anything worth doing is worth overdoing.*

— Mick Jagger

## Data Generation



```
lm.fit()
```



# Contents

1 Introduction

2 Benchmarks

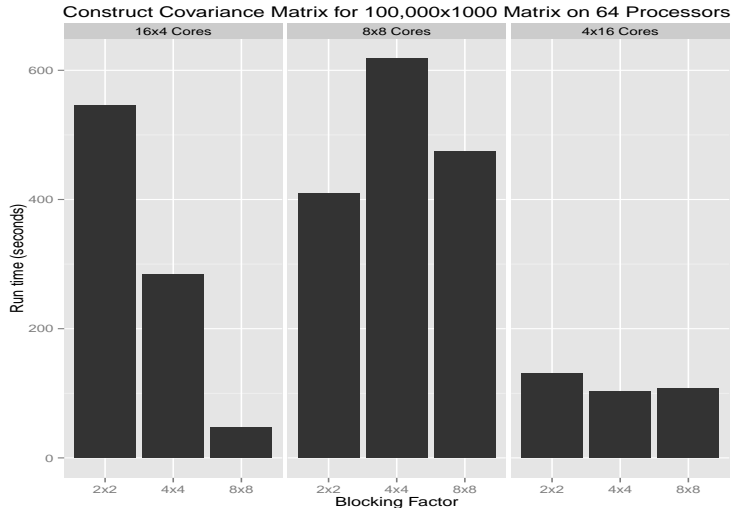
3 Challenges

## Challenges

- Perceptions.
  - “*R? Isn't that slow?*” – HPC people
  - “*HPC? Isn't that hard?*” – R people
- Package loading.
- Profiling.
- Data distribution and performance.



## Covariance Revisited: Distributed Data Parameter Calibration



Thanks for coming!

Questions?

<http://r-pbd.org/>